Towards Trustworthy IoT Network Security: Interpretable Open-Set Intrusion Detection

Ngari Crisphine Macharia and Ning Weng

***Abstract*—Machine learning offers a promising way to detect cyberattacks in Internet of Things (IoT) networks. However, practical deployment is hindered when operators cannot understand model decisions, and unfamiliar attacks are misclassified as known types. We investigate whether interpretable models can address both challenges. We compare Concept Bottleneck Models (CBMs), which predict through human-defined traffic concepts, with NeSy-NIDS, which learns thresholds for expert rules and combines their outputs with a neural network. On random splits of CTU-IoT-23 and CIC-IoT-2023, JointCBM and NeSy-NIDS achieve a weighted F1 within 0.002 of a standard neural network, averaged over five runs. On later CTU-IoT-23 traffic every model loses one attack class, while JointCBM, HybridCBM and NeSy-NIDS still track the baseline on the other classes. The readable concepts and rules only partly explain predictions. Replacing a CBM's learned concept values with values derived directly from declared concept definitions reduces F1 from 0.933 to 0.457-0.492 on CTU-IoT-23. NeSy-NIDS agrees with its neural component on 97% to 100% of flows, but with its rule component on only 32% to 55%. Out-of-distribution detection similarly resists simple explanation: across CTU-IoT-23, CIC-IoT-2023, TON_IoT, six CBM configurations, and five training runs, encouraging closer concept alignment has no consistent effect on AUROC. A CBM whose concept layer is trained before its classifier loses accuracy but detects unfamiliar CTU-IoT-23 attacks as well as the unconstrained network. Seven measurements derived from familiar traffic associate consistently with detection performance on CIC-IoT-2023 and TON_IoT, but on CTU-IoT-23 their association is set by a single model configuration and disappears without it, so none is validated as a predictor. A method for selecting the scoring method from familiar traffic picked the wrong scorer for every held-out family on CIC-IoT-2023. These findings highlight the need to evaluate interpretable intrusion detectors across multiple training runs, architectures, and datasets, with unfamiliar attack families excluded from feature selection and correlations checked for dependence on single configurations.**

***Index Terms*—Concept bottleneck models, Internet of Things, network intrusion detection, open-set recognition, out-of-distribution detection.**

I. Introduction

Internet of Things (IoT) devices make up a large share of network endpoints, and their irregular updates and heterogeneous firmware make them attractive targets for botnet recruitment, DDoS amplification, and command-and-control (C&C) infrastructure. Because IoT devices cannot run endpoint agents, detection has to work from flows captured at a gateway or router tap [1]. The attack mix has recognizable signatures: Mirai-family botnets send high-rate UDP and TCP floods [2], C&C heartbeats beacon at regular inter-arrival times, and reconnaissance scans produce asymmetric packet counts. Such signatures lend themselves to rule- and concept-based representations. At the same time, new Mirai sub-families and DDoS variants keep appearing, so a model trained on a handful of attack families will encounter structurally similar but unseen variants in deployment.

Deep-learning NIDS classify known attacks accurately but raise two problems in deployment. The first is *opacity*: without knowing which traffic features drove a decision, operators cannot validate alerts, adjust policies, or meet regulatory requirements [3], [4]. Post-hoc explainers such as xNIDS [5] give more faithful explanations than LIME or SHAP, but they still explain a black box after the fact. The second is the *open-set problem*: a NIDS trained on a closed set of attacks meets new families in deployment [4] and confidently assigns them to known classes, and common remedies such as retraining or a separate anomaly detector do not use the model's own representation for open-set scoring [6].

In this paper, we study both problems together by building interpretability into the architecture and measuring what that does to open-set detection. The evaluation framework is evaluated (Fig. 1). In **CBMs** [7], predictions pass through a layer of human-defined flow concepts that can be inspected and overridden at test time. In **NeSy-NIDS**, expert attack signatures become threshold rules whose thresholds are learned through k-annealed sigmoid gates; a learned weight blends the rule head with a neural fallback, and a binarized rule vector is kept for audit and OOD scoring.

Contributions.

1) **Accuracy without full fidelity.** Both interpretable architectures match an unconstrained MLP on known classes, but their declared concepts and rules only partly account for their predictions.
2) **Geometry is associated with detectability, not a validated predictor.** Concept-fidelity regularisation has no monotone effect on OOD detection, and known-class geometry tracks OOD AUROC in a consistent direction on two of three corpora; on the third the association rests on a single configuration, so none of the statistics is yet a predictor.
3) **Scorer selection remains open.** The best OOD scorer depends on both dataset and architecture, and a leave-

N. C. Macharia and N. Weng are with the School of Electrical, Computer and Biomedical Engineering, Southern Illinois University Carbondale, IL 62901 USA (e-mail: c.ngari@siu.edu, nweng@siu.edu). Corresponding author: N. C. Macharia.

Code and pre-trained models will be made publicly available upon acceptance, including the full concept and rule definitions with initial thresholds, per-dataset feature lists, and the train/validation/test split and evaluation-subset indices and scripts.

family-out procedure selects it on one corpus but picks the wrong scorer for every held-out family on the other.

Scope. The building blocks are established: concept bottlenecks [7], sigmoid-relaxed threshold rules [8], [9], per-class Mahalanobis scoring [10], and neural/rule hybridisation. We claim no new state-of-the-art score. The contribution is empirical and methodological: a controlled characterisation of where interpretable open-set IoT NIDS hold up and where they do not, with evaluation guidance in the spirit of [4]. We make no streaming, concept-drift, or false-alert-rate claims.

II. Related Work

Post-hoc explanation is the dominant answer to NIDS opacity: xNIDS [5] uses sparse group lasso to produce human-readable firewall rules, and GEAD [11] extracts global rules from a trained deep model. The explanation lives outside the model and can stop matching it when the model changes, so we build interpretability into the architecture instead.

CBMs [7] force predictions through human-specified concepts. Concept embedding models [12] extend concepts to higher-dimensional embeddings, and a trained network can be converted into a CBM post hoc [13]. Concept bottlenecks are known to leak: Mahinpei et al. [14] show that concept models carry predictive information their labels do not capture. Our fidelity tests reproduce this on IoT flows; what we add is its relation to the rule-branch dependence of a parallel neuro-symbolic model and to open-set detectability, under one multi-seed, multi-architecture, multi-dataset protocol. Closest to our setting, Shoukat et al. [15] apply a concept bottleneck with prototype reasoning to intrusion detection in SDN-based industrial networks; we differ in evaluating open-set detection on IoT flows, comparing against a neuro-symbolic rule model, and measuring decision fidelity and operating-threshold behaviour.

Petersen et al. [8] train binary logic-gate networks through smooth relaxations, and Yang et al. [9] inject logical constraints with the straight-through estimator [16]. NeSy-NIDS applies this line of work to IoT flows, where rules are threshold conjunctions over flow features that an analyst can read and override.

The Mahalanobis detector [10] fits class-conditional Gaussians to learned features, and Fort et al. [17] found that feature quality matters more for OOD performance than the scoring rule. Softmax-based alternatives include maximum softmax probability [18] and the energy score [19]. Vaze et al. [20] report that on visual-recognition benchmarks a better closed-set classifier tends to give better open-set recognition; we test whether this empirical relationship transfers to IoT flows. In security, OWAD [6] detects normality shift and CADE [21] isolates drifting samples, both at the system level, whereas we score individual flows. N-BaIoT [2] and Kitsune [1] established per-device anomaly detection for IoT but neither classifies attacks nor explains its decisions. Arp et al. [4] and Apruzzese et al. [22] identify closed-world evaluation and opacity as obstacles to deployment, and De Keersmaeker et al. [23] survey the public IoT security datasets.

TABLE I: Data and evaluation protocol. Split: stratified 70/15/15 per known class, random at the flow level, fixed split seed.

	CTU-IoT-23	CIC-IoT-2023	TON_IoT
Known classes	4	5	5
Held-out families	9	29	5
Features	31 (raw)	39 (z-scored)	23 (18 z-scored)
Known flows per class	23,000	50,000	50,000
Train / val / test (known)	64,400 / 13,800 / 13,800	175k / 37.5k / 37.5k	175k / 37.5k / 37.5k
Unknown pool	32,344	40.2M	373,857
CBM unknown draw	10,000 random	$\leq$2,000/family	$\leq$2,000/family
NeSy unknown draw	10,000 random	344/family	–
AUROC evaluation	5,000 known test vs. 5,000 unknown flows		
Seeds	0–4 (initialisation and batch order only)		

III. Methodology

Fig. 1 shows the evaluation framework, and Algorithms 1 and 2 summarise their training and scoring.

A. Datasets and Protocol

CTU-IoT-23 [24] provides labelled flows from IoT devices captured in a laboratory. Each Zeek conn-log record is joined to per-flow statistics extracted from the packet captures; we use 31 of the resulting features (28 numeric, 3 categorical), selected on known-class training flows only. **CIC-IoT-2023** [25] covers 33 attack scenarios across 105 devices; we use 39 of its flow features. **TON_IoT** [26] is an independent testbed whose network subset is derived from Zeek logs; we use 23 features from its connection and protocol logs. TON_IoT enters only the CBM concept-space analysis (Section IV-D), as a third corpus. Its flow logs carry no per-flag counts or inter-arrival times, so three CTU concepts that need them are replaced by connection-state, responsiveness and web-service concepts.

Table I summarises the protocol. Known classes are class-balanced by random sampling, and every held-out family is excluded from training and validation. The split and the OOD evaluation subset are the same for all seeds.

The known sets are: CTU {Benign, C&C-HeartBeat, DDoS, Okiru}; CIC {Benign, DDoS-SYN_Flood, Mirai-greeth_flood, Recon-PortScan, VulnerabilityScan}; TON {normal, scanning, DDoS, password, injection}. The held-out families are, for CTU, C&C, Attack, C&C-HeartBeat-Attack, C&C-FileDownload, C&C-Torii, C&C-Mirai, C&C-HeartBeat-FileDownload, FileDownload and Okiru-Attack; for CIC, the dataset's other 29 attack classes (DDoS, DoS, Mirai, reconnaissance, spoofing, web-attack, brute-force and backdoor classes); and for TON, DoS, XSS, backdoor, ransomware and MITM. Each known set is a benign baseline plus one representative family from each dominant malicious category in the dataset (e.g. volumetric DDoS, botnet C&C or Mirai-family traffic, and scanning where the dataset has it; TON adds credential and web attacks), the traffic an operator could plausibly label early in a deployment. Most held-out families are sub-families and protocol variants of these categories, which is the operational open-set case.

Preprocessing. CIC and TON numeric features are standardised with a scaler fit on the training split. CTU features enter in raw units (missing values set to -1), with categorical fields integer-encoded from vocabularies built on known-class flows. Concept and rule thresholds were written from documented

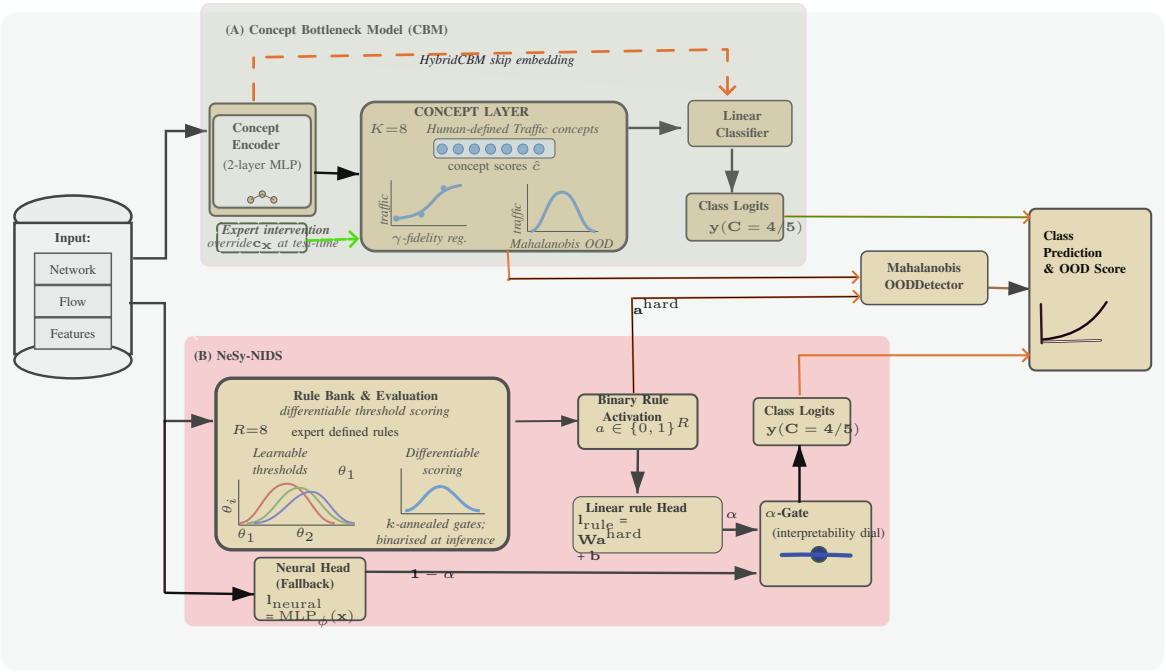


Fig. 1: The Evaluation framework. (A) CBM: concept encoder, concept layer, linear classifier; Mahalanobis OOD scoring in concept space. (B) NeSy-NIDS: rule bank with learnable thresholds and a parallel neural fallback, blended by the learned rule-head weight α. The classifier uses the soft rule activations $\mathbf{a}^{\text{soft}}$; the binarised rule vector $\mathbf{a}^{\text{hard}}$ (Eq. 5) feeds OOD scoring and audit only.

attack signatures and known-class statistics. No preprocessing statistic or concept or rule threshold uses a held-out family. The submitted CTU list (37 features) was an exception: it was selected by random-forest importance on a sample that included every CTU label, the held-out families among them, which may favour CTU OOD separability (CIC and TON use fixed lists). All CTU results here therefore use the same selection re-run on known-class training rows only (31 features; NeSy-NIDS adds back the two counters its expert rules read), and the submitted list is kept as an ablation. The re-selection leaves the accuracy and open-set findings in place, and it raises rather than lowers the AUROC of the two configurations that are most sensitive to the feature list (Supplementary Section SVII, Table S6). Every split and evaluation subset is drawn deterministically by the scripts we release.

Robustness splits. We repeat the CBM evaluation on an alternative known set per dataset (CTU {Benign, DDoS, Okiru, C&C}, promoting C&C to known and demoting C&C-HeartBeat to unknown; CIC {Benign, DDoS-ICMP_Flood, Mirai-udpplain, Recon-PortScan, VulnerabilityScan}) and on a timestamp-ordered CTU split, in which each known class is ordered by capture time and the earliest 70% trains and the latest 15% tests. For CTU, the feature list is re-selected by the same random-forest procedure on each experiment's own known-class training rows, so the family it holds out (and any test-period flow) never informs its features. Each of these splits draws its own class-balanced sample; a control that uses the main split's 31-feature list on the same sample separates the effect of re-selection from that of the sample, and the accuracy and open-set conclusions hold under both.

B. Concept Bottleneck Models

A CBM [7] factors the classifier into two stages,

$$\hat{\mathbf{c}} = g(\mathbf{x}), \quad \hat{y} = h(\hat{\mathbf{c}}), \tag{1}$$

where $g : \mathbb{R}^d \to [0,1]^K$ is the concept encoder (an MLP $d{\to}256{\to}256{\to}64{\to}K$ with LayerNorm and ReLU) and $h : [0,1]^K \to \mathbb{R}^C$ is a linear classifier over the C known classes. We define $K{=}8$ concepts per dataset (e.g. *is_high_rate*: Rate$>$1000 pps; *is_beaconing*: 1.5$<$IAT$<$5 s $\wedge$ 0.1$<$Rate$<$5; *is_incomplete_handshake*: syn_count$>$2 $\wedge$ fin_count$=$0), each a binary label produced by a threshold rule on the raw features (Table II). Because the labels are deterministic functions of the input, *concept accuracy* measures how faithfully the bottleneck reproduces these declared predicates, not whether the predicates are semantically correct. $K{=}8$ balances vocabulary against the sample size the shared-covariance Mahalanobis fit needs.

Training variants. *JointCBM* minimises

$$\mathcal{L} = \mathcal{L}_{\text{CE}} + \lambda\, \mathcal{L}_{\text{BCE}} + \frac{\gamma}{K} \|\hat{\mathbf{c}} - \mathbf{c}_{\text{gt}}\|^2, \tag{2}$$

where $\mathbf{c}_{\text{gt}} \in \{0,1\}^K$ are the concept labels, λ weights concept supervision, and $\gamma \in \{0, 0.1, 0.5, 1.0\}$ weights an additional squared-error concept-fidelity term. Both terms push the predicted concepts toward the same labels. *SequentialCBM* trains and freezes g on the concept labels before training h. *HybridCBM* concatenates the 64-dimensional encoder embedding with the K concepts as classifier input, so the classifier can use information outside the concept vocabulary. *Post-hoc CBM* fits a logistic-regression probe per concept to the frozen MLP embedding and a logistic-regression classifier to the probe outputs, both on the training split.

TABLE II: Concept definitions: the threshold rule that produces each binary concept label. CTU thresholds are in raw units; CIC thresholds are z-scores with raw-unit equivalents in parentheses; TON_IoT uses raw Zeek fields and replaces three CTU concepts that need fields its logs lack.

Concept	Definition
CTU-IoT-23 (raw units)	
is_beaconing	1.5<IAT<5.0s ∧ 0.1<Rate<5.0
is_high_rate	Rate >1000 pps
is_syn_heavy	syn_count >8
is_short_connection	0≤duration<0.01s
is_persistent	duration >2.5s
is_asymmetric	resp_pkts=0 ∧ orig_pkts>1
is_single_packet_probe	orig_pkts=1
is_incomplete_handshake	syn_count>2 ∧ fin_count=0
CIC-IoT-2023 (z-scored; raw-unit equivalent in parentheses)	
is_high_rate	rate >0.15 (>2.4×10^4 pps)
is_syn_flood	syn_count >1.5 (>79)
is_udp_dominant	tcp <−1.0 (<0.26)
is_short_connection	iat <−0.00316 (<4.7×10^{-4} s)
is_large_payload	avg >0.4 (>529 B)
is_port_scan	rst_count >1.5 (>4.8)
is_high_variance	variance >0.5 (>8.2×10^5)
is_persistent	fin_count >1.0 (>0.78)
TON_IoT (raw Zeek conn-log units)	
is_high_rate	(src_pkts+dst_pkts)/duration >1000
is_short_connection	0≤duration<0.01s
is_persistent	duration >2.5s
is_asymmetric	dst_pkts=0 ∧ src_pkts>1
is_single_packet_probe	src_pkts=1
is_incomplete_handshake	conn_state ∈ {S0,S1,REJ,RSTOS0,RSTRH,SH,SHR}
is_responsive	dst_pkts >0
is_web_service	service ∈ {http, ssl}

Algorithm 1 CBM Training and OOD Scoring

Require: $\{(\mathbf{x}_i, \mathbf{c}_i, y_i)\}$, λ, γ
1: Initialise encoder g (d→256→256→64→K, LayerNorm+ReLU), linear classifier h
2: **for** each epoch and mini-batch $(\mathbf{X}, \mathbf{C}, \mathbf{Y})$ **do**
3: $\hat{\mathbf{C}}=g(\mathbf{X})$; $\mathcal{L}=\mathcal{L}_{\mathrm{CE}}(h(\hat{\mathbf{C}}), \mathbf{Y})+\lambda\mathcal{L}_{\mathrm{BCE}}(\hat{\mathbf{C}}, \mathbf{C})+\frac{\gamma}{K}\|\hat{\mathbf{C}}-\mathbf{C}\|^2$; update g, h
4: **end for**
5: $\boldsymbol{\mu}_c$=mean$\{g(\mathbf{x}_i): y_i=c\}$; $\boldsymbol{\Sigma}$= regularised shared covariance
6: **Inference:** $\hat{\mathbf{c}}=g(\mathbf{x})$; $\hat{y}=\arg\max h(\hat{\mathbf{c}})$; $s(\mathbf{x})$ by Eq. 3

OOD scoring. We apply per-class Mahalanobis scoring [10] in concept space. Class means $\boldsymbol{\mu}_c$ and a shared covariance $\boldsymbol{\Sigma}$ ($\varepsilon=10^{-4}$ ridge, pseudoinverse) are fitted to training concept vectors, and the anomaly score is

$$s(\mathbf{x}) = \min_c (\hat{\mathbf{c}} - \boldsymbol{\mu}_c)^\top \boldsymbol{\Sigma}^{-1} (\hat{\mathbf{c}} - \boldsymbol{\mu}_c). \tag{3}$$

C. NeSy-NIDS: Differentiable Rule Learning

Rule templates. The $R=8$ rule templates per dataset are written from domain knowledge. Each is a conjunction of threshold conditions; representative CTU-IoT-23 rules are `ddos_short_burst` (Rate>1000 ∧ duration<0.1; DDoS), `okiru_low_rate` (Rate<500 ∧ duration<2.5; Okiru) and `benign_single_probe` (orig_pkts<2 ∧ Rate<100; Benign), and several rules start from concept thresholds (e.g. `cc_heartbeat_beaconing` from *is_beaconing*). The full banks with initial thresholds and declared target classes are in Supplementary Table S8. Let $\mathcal{I}_r^{>}$ and $\mathcal{I}_r^{<}$ index the greater-than and less-than conditions of rule r. Following [8], each indicator is replaced by a sigmoid gate and the conjunction by a product:

$$a_r(\mathbf{x}; k) = \prod_{i\in\mathcal{I}_r^{>}} \sigma\big(k(x_i-\theta_i)\big) \cdot \prod_{j\in\mathcal{I}_r^{<}} \sigma\big(k(\theta_j-x_j)\big), \tag{4}$$

so the thresholds θ receive gradients throughout training.

Annealing and binarisation. The steepness k rises linearly from $k_0=1$ to $k_T=10$ over $T_{\mathrm{warm}}=30$ epochs, $k(t) = k_0 + \min(t/T_{\mathrm{warm}}, 1)(k_T-k_0)$. Training and classification use the soft activations of Eq. 4. For audit and OOD scoring, each condition is thresholded and the rule is their logical conjunction,

$$a_r^{\mathrm{hard}} = \prod_{i\in\mathcal{I}_r^{>}} \mathbf{1}[x_i>\theta_i] \prod_{j\in\mathcal{I}_r^{<}} \mathbf{1}[x_j<\theta_j], \tag{5}$$

which is not the same operation as thresholding the soft product.

Architecture and rule-head weight. A linear rule head maps the soft activations, scaled by learned per-rule importance weights, to rule logits $\mathbf{l}_{\mathrm{rule}}$. A parallel MLP (hidden layers of 128 and 64 units, LayerNorm and ReLU) on the raw features produces neural logits $\mathbf{l}_{\mathrm{neural}}$, and

$$\mathbf{l} = \alpha\, \mathbf{l}_{\mathrm{rule}} + (1-\alpha)\, \mathbf{l}_{\mathrm{neural}}, \qquad \alpha=\sigma(w_\alpha), \tag{6}$$

with w_α learned jointly. An optional penalty $\lambda_\alpha(1-\alpha)^2$ pushes α upward. α is a single global mixing weight; because the two heads can emit logits of different magnitudes, its value alone does not measure how much the prediction relies on the rules (Section IV-C).

OOD scoring in rule space. The binary vector $\mathbf{a}^{\mathrm{hard}}(\mathbf{x}) \in \{0,1\}^R$ is scored with Eq. 3, fitted on known-class training activations. Binary vectors are not Gaussian, so this is an approximation; Section IV-B compares it with binary-native and logit-space scorers.

Algorithm 2 NeSy-NIDS Training and OOD Scoring

Require: $\{(\mathbf{x}_i, y_i)\}$, templates with initial thresholds $\boldsymbol{\theta}^{(0)}$, λ_α, $k_0=1$, $k_T=10$, T_{warm}
1: Initialise $\boldsymbol{\theta}=\boldsymbol{\theta}^{(0)}$, rule head $(\mathbf{W}, \mathbf{b}, \mathbf{w}_{\mathrm{imp}})$, MLP ϕ, gate w_α
2: **for** epoch $t = 1, \ldots, T$ **do**
3: $k=k_0+\min(t/T_{\mathrm{warm}}, 1)(k_T-k_0)$
4: **for** each mini-batch $(\mathbf{X}, \mathbf{Y})$ **do**
5: $\mathbf{a}=\mathbf{a}(\mathbf{X}; k)$ (Eq. 4, soft); $\mathbf{l}=\alpha\, \mathbf{l}_{\mathrm{rule}}(\mathbf{a})+(1-\alpha)\,\mathrm{MLP}_\phi(\mathbf{X})$, $\alpha=\sigma(w_\alpha)$
6: $\mathcal{L}=\mathcal{L}_{\mathrm{CE}}(\mathbf{l}, \mathbf{Y})+\lambda_\alpha(1-\alpha)^2$; update $\boldsymbol{\theta}, \mathbf{W}, \mathbf{b}, \mathbf{w}_{\mathrm{imp}}, \phi, w_\alpha$
7: **end for**
8: **end for**
9: Fit $\boldsymbol{\mu}_c$, $\boldsymbol{\Sigma}$ on $\mathbf{a}^{\mathrm{hard}}(\mathbf{x}_i)$ (Eq. 5) of the training flows
10: **Inference:** classify $\hat{y}=\arg\max \mathbf{l}$ from soft $\mathbf{a}(\mathbf{x}; k_T)$; score $s(\mathbf{x})$ by Eq. 3 on $\mathbf{a}^{\mathrm{hard}}(\mathbf{x})$ (OOD and audit)

D. Open-Set Decision Rule

Both classifiers are closed-set, so we couple the anomaly score to the classifier through a rejection threshold τ:

$$\hat{y}_{\mathrm{open}}(\mathbf{x}) = \begin{cases} \textit{unknown}, & s(\mathbf{x}) > \tau, \\ \arg\max_c\ l_c(\mathbf{x}), & \text{otherwise}, \end{cases} \tag{7}$$

where l_c is the c-th class logit. τ is the 95th percentile of scores on the known-class validation split, so about 5% of known traffic is rejected; no unknown flow is used to set it.

E. Baselines, Metrics and Training

The MLP baseline has the same encoder as the CBM followed by a linear classifier and is scored with Mahalanobis on its 64-dimensional embedding. A decision tree (DT; depth chosen from $\{5, 10, 15, \infty\}$ by validation F1) is paired with per-class Mahalanobis on the raw training features, and a random forest (RF [27]; 100 trees, depth 10) is scored with $1-\max(\hat{p})$. We report weighted F1 on the known test split, and AUROC and TPR at 5% FPR with known test flows as negatives and held-out unknown flows as positives. All hyperparameters and early stopping use the known-class validation split. Adam (lr=10^{-3}, batch 512); CBMs train for up to 50 epochs (patience 10) with λ=0.5, chosen from $\{0.1, 0.5, 1.0\}$ by known-class validation F1, which varies by at most 0.002 across that grid; on the original-feature splits, raising λ from 0.1 to 1.0 raises concept accuracy on both datasets but moves OOD AUROC in opposite directions (CTU 0.790→0.884, CIC 0.752→0.724, five seeds; Supplementary Table S4); NeSy-NIDS trains for up to 60 epochs (patience 12) with $\lambda_\alpha \in \{0, 0.1\}$.

Evaluation subsets. Each analysis draws one evaluation subset and reuses it for every model and seed it compares; the NeSy scorer ablation (Table VI) scores all five scorers on a subset drawn per seed. Different analyses draw independently under the same 5,000/5,000 protocol, so the same model can differ slightly between tables. For the CIC MLP, Table III gives 0.807±.007 (five seeds), the scorer × architecture analysis 0.807±.008 (five seeds, another subset), and the post-hoc detector comparison 0.815 (the seed-0 model on a third subset). These subsets span 0.807–0.815, a range of 0.008 that is below the cross-architecture and cross-scorer gaps we interpret (≥ 0.05), and we compare values only within one analysis.

IV. Results

A. Classification Accuracy

Table III shows NeSy-NIDS and every JointCBM variant tracking the MLP on both datasets. Paired per-seed F1 differences from the MLP have means of at most 0.0001 in magnitude on CTU-IoT-23 (no seed beyond 0.0004) and −0.0004 to +0.0002 on CIC-IoT-2023 (no seed beyond 0.0026). These are small observed differences, well inside seed variation, not a demonstration of statistical equivalence. For JointCBM (γ=0) the match also holds per class: every class is within 0.002 of the MLP. The weighted F1 hides a spread across classes: the weakest are CTU Okiru (0.86) and DDoS (0.90), and CIC Recon-PortScan (0.64), VulnerabilityScan (0.70) and Benign (0.76).

The rule-only head reaches only 0.566 on CTU-IoT-23 and 0.750 on CIC-IoT-2023, so the neural fallback covers traffic the eight rule templates do not. SequentialCBM and Post-hoc CBM lose much of their accuracy (CTU 0.567/0.576; CIC 0.732/0.746): this is consistent with a linear classifier over concepts learned separately from the class labels being unable to recover the class boundary. The γ variants have near-identical F1 in the main runs; one heavily regularised run in the supplementary (λ, γ) grid (λ=1.0, γ=0.5, seed 0) fell to F1 0.748.

Robustness. On the alternative known sets the MLP reaches F1 0.935 (CTU) and 0.820 (CIC), and every JointCBM variant matches it within 0.001 on both. The timestamp-ordered split lowers four-class weighted F1 from 0.933 to 0.646–0.648 for every model except SequentialCBM (0.330; Table IV). The CTU *DDoS* label spans two capture regimes (a high-rate flood, then a near-zero-rate capture), so the temporal split trains on one and tests on the other, and every model scores F1 0.00 on DDoS on every seed; on the same sample with the main split's 31-feature list DDoS also collapses for every model and seed except two JointCBM γ=0.5 seeds, so the cause is the label's temporal structure, not feature selection. On the three temporally-coherent families the interpretable models match the MLP in aggregate (Coh-3 0.862–0.863) while trading per-class strengths (JointCBM γ=0: Benign 0.841 vs. 0.958, Okiru 0.961 vs. 0.844); only SequentialCBM, which loses Benign, falls behind. The accuracy match is therefore a matched-distribution result.

B. Open-Set Detection

CTU-IoT-23. Under the leak-free feature selection the highest observed means come from NeSy-NIDS with the α penalty (0.909±.014) and NeSy-NIDS (0.904±.014), followed by the MLP (0.874±.002), SequentialCBM (0.872±.026) and Post-hoc CBM (0.871±.024); the NeSy-MLP gap is about two seed std. The four JointCBM variants reach 0.81–0.84 (seed std 0.04–0.05), HybridCBM 0.822, and the DT and RF baselines 0.679 and 0.657. SequentialCBM, whose concept head is trained on the concept labels alone and frozen, detects as well as the MLP while its closed-set F1 is only 0.567: on CTU a sequentially trained bottleneck gives up accuracy but not open-set detection. With the originally submitted feature list on the same rows, SequentialCBM and HybridCBM reach only 0.782 and 0.772, so known-class-only selection raises them (ablation, Supplementary Table S6).

CIC-IoT-2023. The MLP (0.807±.007) and the DT's raw-feature Mahalanobis (0.802) have the highest AUROC. The JointCBM variants reach 0.699–0.717; SequentialCBM and HybridCBM reach 0.667 and 0.669, and NeSy-NIDS with Mahalanobis only 0.657 (0.691 with the α penalty). Unlike on CTU, SequentialCBM is among the weakest detectors here. The low-dimensional concept and rule spaces separate these unknowns less well than the 64-dimensional MLP embedding and 39 raw features, though we ran no dimension-matched control and do not attribute the gap to dimensionality alone.

Concept-fidelity regularisation is not a lever. The sign of $\Delta=\gamma_{1.0}-\gamma_0$ is opposite on the two main splits, and across five dataset/split combinations it is positive on two and negative on three, with $|\Delta| \leq 0.06$ (Table V). In a joint (λ, γ) grid (three seeds per cell; Supplementary Section SXIV) the sign of the γ effect also depends on λ. An effect of this size can be read off a single configuration and a single seed, so claims about it need multi-seed replication under a fixed scoring protocol.

Which scorer works depends on the representation. Table VI compares per-class Mahalanobis on the NeSy rule vector with two binary-native distances to the nearest known-class centroid (Hamming, Jaccard), the logit-space energy score

TABLE III: Weighted F1 on the known-class test split; OOD AUROC and TPR@5%FPR (known vs. held-out unknown test flows). F1 is the five-seed mean (seed std $\leq$0.0015; SequentialCBM, Post-hoc CBM and rule-only $\leq$0.0152); AUROC and TPR show mean$\pm$std over five seeds. **Bold** marks the highest observed mean in each OOD column, not an established ranking. OOD scores are per-class Mahalanobis fitted on training vectors: activations for the CBM, NeSy and MLP rows, concept-probe outputs for Post-hoc CBM, and ‡raw features for DT (this score does not use the tree, so it has no seed variance); †RF uses $1-\max \hat{p}$. CBM, MLP, Post-hoc CBM, DT and RF rows share the same evaluation flows. CTU-IoT-23 uses the known-class-only feature selection (leak-free; Section III-A); the originally submitted feature list, run on the same rows, is reported as an ablation in Supplementary Table S6.

	CTU-IoT-23			CIC-IoT-2023		
Model	**F1**	**AUROC**	**TPR@5%**	**F1**	**AUROC**	**TPR@5%**
MLP Baseline	0.9331	0.874±.002	0.519±.001	0.8186	**0.807±.007**	0.516±.044
JointCBM (γ=0)	0.9330	0.823±.044	0.376±.199	0.8183	0.717±.017	0.361±.029
JointCBM (γ=0.1)	0.9331	0.805±.035	0.335±.116	0.8188	0.700±.009	0.344±.035
JointCBM (γ=0.5)	0.9330	0.844±.036	0.475±.152	0.8184	0.702±.007	0.333±.031
JointCBM (γ=1.0)	0.9331	0.834±.050	0.520±.088	0.8182	0.699±.013	0.352±.028
SequentialCBM	0.5673	0.872±.026	0.550±.049	0.7321	0.667±.005	0.177±.015
HybridCBM	0.9331	0.822±.055	0.461±.049	0.8188	0.669±.015	0.173±.031
Post-hoc CBM	0.5758	0.871±.024	0.462±.100	0.7457	0.652±.018	0.234±.029
NeSy-NIDS	0.9331	0.904±.014	**0.586±.098**	0.8186	0.657±.028	0.252±.056
NeSy + α-reg (λ_α=0.1)	0.9330	**0.909±.014**	0.575±.101	0.8198	0.691±.018	0.253±.081
Rule-only	0.5660	0.870±.011	0.427±.022	0.7498	0.668±.038	0.189±.066
DT‡	0.9337	0.679±.000	0.251±.000	0.8145	0.802±.000	**0.549±.000**
RF†	0.9338	0.657±.034	0.446±.003	0.8057	0.704±.006	0.181±.005

TABLE IV: Temporal (timestamp-ordered) split on CTU-IoT-23, features re-selected on the temporal-train rows (Section III-A): per-class known F1, four-class weighted F1 ($F1_4$), weighted F1 over the three temporally coherent families (Coh-3; DDoS excluded), and OOD AUROC (mean$\pm$std over five seeds). DDoS F1 is 0.00 for every model and seed. $F1_4$ and Coh-3 seed std $\leq$0.001 except SequentialCBM (0.03, 0.04). Each known class has a 3,451-flow test support.

Model	Benign	C&C-HB	DDoS	Okiru	$F1_4$	Coh-3	AUROC
MLP baseline	0.958	0.789	0.00	0.844	0.647	0.863	0.746±.010
JointCBM γ=0	0.841	0.788	0.00	0.961	0.647	0.863	0.758±.062
JointCBM γ=0.1	0.840	0.788	0.00	0.961	0.647	0.863	0.783±.040
JointCBM γ=0.5	0.840	0.788	0.00	0.961	0.647	0.863	0.765±.044
JointCBM γ=1.0	0.880	0.789	0.00	0.922	0.648	0.863	0.764±.050
SequentialCBM	0.104	0.773	0.00	0.445	0.330	0.441	0.839±.012
HybridCBM	0.879	0.787	0.00	0.922	0.647	0.862	0.716±.059
NeSy-NIDS	0.917	0.786	0.00	0.883	0.646	0.862	0.888±.025

TABLE V: JointCBM OOD AUROC vs. γ (mean over five seeds) on every split; Δ=$\gamma_{1.0}-\gamma_0$.

Split	Data	γ=0	γ=0.1	γ=0.5	γ=1.0	Δ
Standard	CTU	0.823	0.805	0.844	0.834	+0.011
Alternative	CTU	0.837	0.837	0.832	0.809	−0.028
Temporal	CTU	0.758	0.783	0.765	0.764	+0.006
Standard	CIC	0.717	0.700	0.702	0.699	−0.018
Alternative	CIC	0.763	0.765	0.765	0.712	−0.051

$-\log \sum_c e^{l_c}$ [19], and a z-normalised Mahalanobis+energy average. On CTU the two rule-vector distances and per-class Mahalanobis all work (0.86–0.90) while energy is below chance (0.46). On CIC no activation-space scorer exceeds 0.66 and only the energy-based scorers detect (0.747–0.763). The Gaussian assumption is therefore not what limits the binary rule

TABLE VI: NeSy-NIDS OOD AUROC by scorer (mean over five seeds; std in parentheses). Mahalanobis, Hamming and Jaccard score the binary rule vector; Energy scores the logits; Comb. is their z-normalised average. Bold: highest observed mean per dataset.

Scorer	Space	CTU-IoT-23	CIC-IoT-2023
Mahalanobis (per-class)	activations	**0.904** (.014)	0.657 (.028)
Hamming to centroid	activations	0.858 (.007)	0.650 (.008)
Jaccard to centroid	activations	0.888 (.007)	0.640 (.015)
Energy	logits	0.461 (.076)	0.747 (.028)
Mahalanobis + Energy	both	0.661 (.105)	**0.763** (.020)

space; whether that space separates known from unknown flows at all is. Running Mahalanobis and energy on the MLP and all six CBM configurations shows that the ranking depends on architecture as well (Table VII). On CTU Mahalanobis is better for every architecture, by 0.23–0.49. On CIC energy is better for the four JointCBM variants (0.732–0.769 vs. 0.702–0.722), while Mahalanobis stays better for the MLP, SequentialCBM and HybridCBM. No scorer is best across both datasets and all architectures.

Scorer selection. Whether the right scorer can be chosen without the novel traffic is a separate question from how well a given scorer performs. We tested one procedure: hold out one known attack family, retrain JointCBM (γ=0) on the rest, and select the scorer (Mahalanobis or energy) that better detects the held-out family (five seeds per family; Table VIII). On CTU the family-majority vote recovers the true-best Mahalanobis (two of three families, 9/15 folds). On CIC it picks Mahalanobis for all four families (15/20 folds) although energy is best on the true unknowns: the held-out known families stay separable

TABLE VII: OOD AUROC of per-class Mahalanobis vs. the energy score by architecture (mean over five seeds; std in parentheses). All models share one evaluation subset, drawn independently of Table III. Bold: better scorer per architecture and dataset.

	CTU-IoT-23		CIC-IoT-2023	
Architecture	Mahal.	Energy	Mahal.	Energy
MLP	**0.879** (.001)	0.386 (.073)	**0.807** (.008)	0.615 (.020)
JointCBM (γ=0)	**0.830** (.044)	0.476 (.111)	0.722 (.014)	**0.734** (.033)
JointCBM (γ=0.1)	**0.812** (.034)	0.494 (.136)	0.705 (.006)	**0.732** (.028)
JointCBM (γ=0.5)	**0.851** (.037)	0.570 (.138)	0.706 (.004)	**0.757** (.014)
JointCBM (γ=1.0)	**0.842** (.049)	0.612 (.124)	0.702 (.014)	**0.769** (.012)
SequentialCBM	**0.875** (.026)	0.604 (.085)	**0.665** (.006)	0.636 (.017)
HybridCBM	**0.827** (.055)	0.365 (.047)	**0.670** (.014)	0.629 (.023)

TABLE VIII: Leave-family-out scorer selection (JointCBM γ=0, five seeds per held-out family). AUROCs are five-fold means; "Winner" counts folds selecting Mahalanobis / energy on the held-out family, so it can disagree with the means.

Data	Held-out family	Mahal. AUROC	Energy AUROC	Winner (M/E)
CTU	C&C-HeartBeat	0.90	0.50	5/0
CTU	DDoS	0.86	0.55	4/1
CTU	Okiru	0.22	0.49	0/5
CIC	DDoS-SYN_Flood	0.86	0.54	5/0
CIC	Mirai-greeth_flood	0.85	0.90	4/1
CIC	Recon-PortScan	0.75	0.72	3/2
CIC	VulnerabilityScan	0.70	0.64	3/2

True-best scorer: Mahalanobis (CTU-IoT-23), energy (CIC-IoT-2023).

in activation space while the true unknowns do not, so the proxy is confidently wrong on the corpus where the scorer must switch.

This procedure therefore did not reliably identify the best scorer.

Operating point. Table IX evaluates the deployed $(C{+}1)$-class detector at τ (Section III-D). The known rejection rate stays near its 5% target (4.2–5.8%), and τ has a coefficient of variation across seeds of at most 0.11. Open-set macro-F1 is moderate on CTU and low on CIC, where unknown recall is low (for NeSy-NIDS, seed 0, the scorer rejects about 19% of the pooled unknown flows).

Because unknown flows outnumber known ones about 17:1 in this pool, the unrejected unknowns mainly lower known-class precision, so macro-F1 depends on prevalence. Per-family AUROC ranges widely and the two architectures catch different families (per-family results for both: Supplementary Section SXII): for JointCBM it runs from near-perfect (CTU C&C-Mirai 0.994; CIC DDoS-PSHACK_Flood 0.993) to below chance (CIC Mirai-greip_flood 0.15, DDoS-UDP_Fragmentation 0.24), so a strong pooled AUROC does not mean the deployed detector flags every novel family at a usable rate. Against five post-hoc detectors on the seed-0 MLP (Table X), the confidence scores (MSP [18], ODIN, energy) sit near or below chance on CTU and are modest on CIC, while the feature-space detectors (deep kNN [28] and Mahalanobis) are strongest; we make no state-of-the-art claim.

TABLE IX: Open-set $(C{+}1)$ evaluation at τ (5% FPR on known validation), mean (std) over five seeds. Known rej.: fraction of known test flows rejected. Macro-F1: over C known classes plus "unknown". Eq.-fam: AUROC and unknown recall at τ averaged over unknown families with equal weight. The JointCBM row is γ=0.

Data	Model	Known rej. (%)	Macro-F1	Eq.-fam AUROC	Eq.-fam recall
CTU	MLP	4.9 (0.2)	0.765 (.002)	0.948	0.88
CTU	JointCBM	5.0 (0.5)	0.654 (.086)	0.818	0.51
CTU	SequentialCBM	5.8 (0.1)	0.489 (.018)	0.746	0.49
CTU	HybridCBM	5.6 (0.1)	0.646 (.018)	0.664	0.35
CTU	NeSy-NIDS	4.8 (0.4)	0.788 (.014)	0.962	0.83
CIC	MLP	5.2 (0.1)	0.267 (.019)	0.806	0.51
CIC	JointCBM	5.2 (0.2)	0.198 (.012)	0.722	0.36
CIC	SequentialCBM	5.0 (0.1)	0.135 (.004)	0.673	0.16
CIC	HybridCBM	5.2 (0.2)	0.131 (.008)	0.676	0.16
CIC	NeSy-NIDS	4.2 (0.6)	0.145 (.017)	0.641	0.20

TABLE X: Post-hoc OOD detectors on the unconstrained MLP (AUROC; seed-0 model on its own evaluation subset).

Dataset	MSP	ODIN	Energy	kNN	Mahal.
CTU-IoT-23	0.281	0.367	0.442	**0.876**	0.869
CIC-IoT-2023	0.619	0.714	0.607	**0.840**	0.815

C. Fidelity of the Concept and Rule Representations

Readable concepts or rules need not account for the predictions made from them, so we test fidelity separately from accuracy (Table XI). **CBM.** Three representations must be kept apart: the soft predicted concepts, their thresholded values, and the declared predicates used as labels. A linear probe on the soft predicted concepts gains little over one on their thresholded values on CTU-IoT-23 (except for two γ=1.0 seeds) and more on CIC-IoT-2023, where the gain grows rather than shrinks with γ (Table XI). A small gain does not establish fidelity, because thresholded predictions can disagree with the declared predicates and still carry class information. We therefore also pass the declared predicates through the model's own classifier, and train a logistic regression on them directly. Both fall far below the learned concepts (on CTU the replacement F1 has seed std 0.02–0.04), so under the evaluated linear classifiers the declared predicates do not reproduce the model's accuracy. Replacement also changes the classifier's input distribution, so these drops test decision fidelity under predicate replacement; they are not a decomposition of accuracy by source. Raising γ improves concept-label agreement on CTU (87.4% to 94.2%; Supplementary Section SII) but leaves the gap within seed variability.

NeSy-NIDS. The hard conjunction almost always agrees with thresholding the soft product, most soft activations are already near 0 or 1 before the cut (known test split), and classifying from the hard activations changes 0.0% (CTU) and 0.9% (CIC) of known test predictions at λ_α=0 (0.0% and up to 3.2% across the sweep), so binarisation has little effect on the blended predictions, a different question from whether the rules determine them. At λ_α=0 the deployed prediction agrees with the neural-only head on 100% and 99% of flows and with the rule-only head on only 32% and 45% (seed std 17 and 9 points); raising λ_α over $\{0, 0.05, 0.1, 0.2, 0.5, 1.0\}$ (each setting trained from scratch) increases α at essentially constant

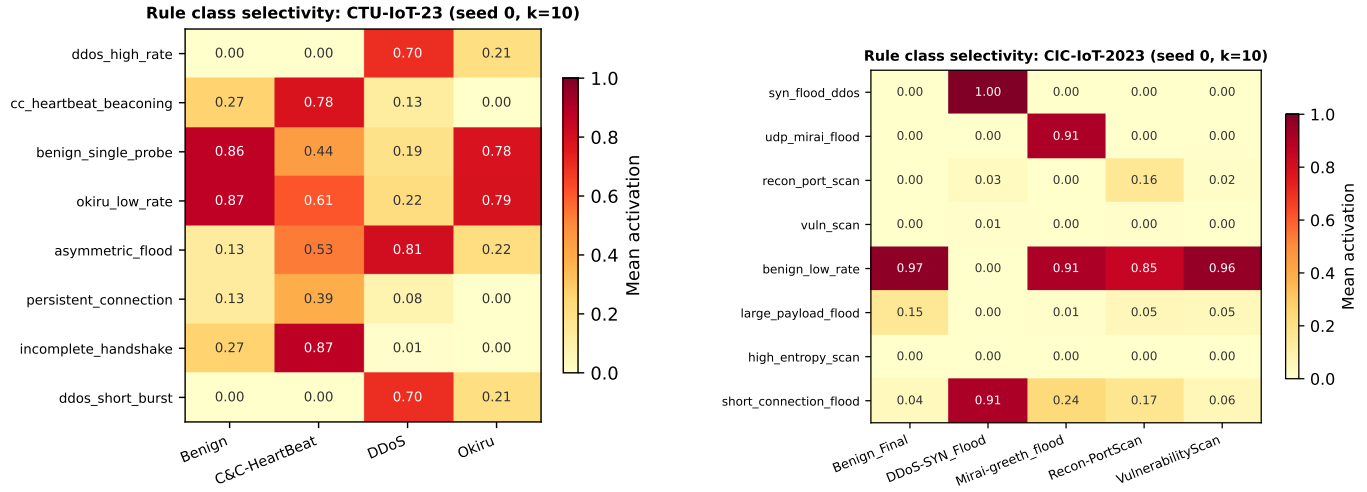


Fig. 2: Mean binarised rule activation (Eq. 5) per known class on the known-class validation split (seed-0 model), CTU-IoT-23 (left) and CIC-IoT-2023 (right). CTU rules fire selectively, leaving within-class variance for the Mahalanobis scorer; CIC rules fire near-deterministically, so there is little to separate on. The pattern is descriptive, not causal (Fig. 4).

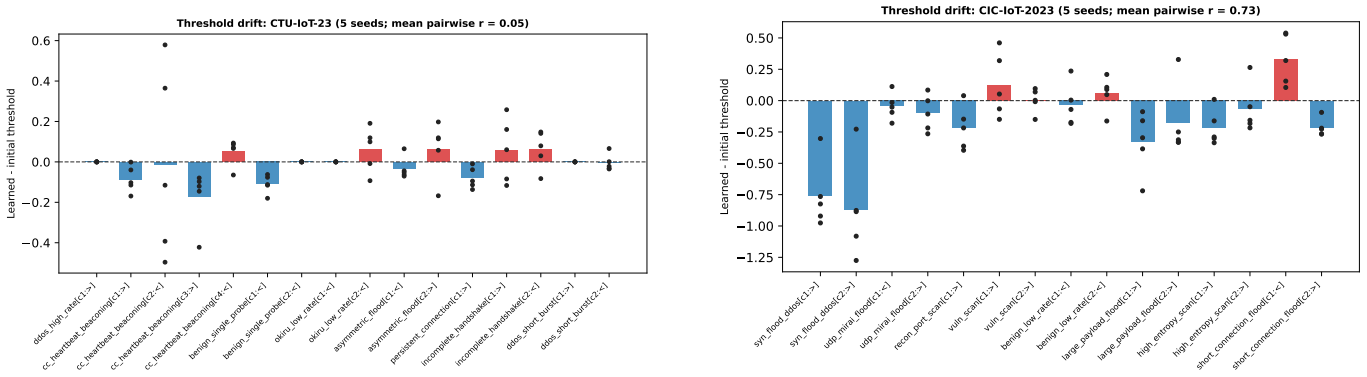


Fig. 3: Per-condition threshold drift (learned − initial) for every NeSy-NIDS rule condition, CTU-IoT-23 (left) and CIC-IoT-2023 (right). Bars are the five-seed mean; dots are individual seeds. Drifts recur across seeds on CIC (mean pairwise r=0.73) but not on CTU (r=0.05), so we read the CIC refinement as systematic and the CTU pattern as indicative only; reproducibility is not semantic validation.

TABLE XI: Fidelity summary (five seeds; per-seed detail in the Supplementary Material). CBM ranges span $\gamma \in \{0, 0.1, 0.5, 1.0\}$; NeSy ranges span $\lambda_\alpha \in \{0, \ldots, 1.0\}$.

	CTU-IoT-23	CIC-IoT-2023
JointCBM		
F1, predicted concepts	0.933	0.818–0.819
F1, declared predicates through own classifier	0.457–0.492	0.624–0.645
F1, logistic regression on declared predicates	0.575	0.718
Probe accuracy, soft minus thresholded concepts	0.001–0.084	0.041–0.058
NeSy-NIDS		
Rule-head weight α	0.38–0.89	0.32–0.95
OOD AUROC across the λ_α sweep	0.893–0.909	0.657–0.691
Prediction agrees with neural-only head	100%	97–99%
Prediction agrees with rule-only head	32–51%	45–55%
Hard conjunction agrees with thresholded soft product	>99.9%	99.7–99.9%
Soft activations within 0.1 of 0 or 1	0.87–0.96	0.75–0.80

known-class validation F1 (about 0.933 CTU, 0.818–0.819 CIC), and rule agreement rises only modestly with it (CTU 32% to 51%, CIC 45–55%), staying well below the neural agreement (Table XI). OOD AUROC stays within 0.016 on CTU and rises by 0.02–0.03 on CIC as α grows; given the branch agreement, which is observational, we do not attribute that rise to greater rule reliance. The rules fire in two regimes (Fig. 2): selectively on CTU-IoT-23, leaving within-class variance for the Mahalanobis scorer, but near-deterministically on CIC-IoT-2023, which coincides with the lower CIC AUROC. The learned thresholds drift reproducibly across seeds on CIC (mean pairwise correlation of the per-condition drifts 0.73) but not on CTU (0.05; Fig. 3); reproducibility is not semantic validation. Rule quality is also uneven (Table XII): precision against each rule's declared target class shows the name often mismatches behaviour (`okiru_low_rate` fires mostly on Benign, `vuln_scan` almost never on VulnerabilityScan, and `benign_low_rate` as often on Mirai as on Benign), and several CIC thresholds vary strongly across seeds. In both architectures the inspectable layer is an honest record of concept or rule activation, but not a sufficient explanation of the final classification.

TABLE XII: Per-rule quality of the NeSy-NIDS rule bank on the known test split (mean over five seeds). Cov.: fraction of flows activating the rule. Prec.: fraction of activating flows in the rule's dominant class. Target prec.: fraction in the rule's declared target class (Table S8), seed std in parentheses. Thr.: relative std of the learned threshold across seeds.

Rule	Cov.	Dominant	Prec.	Target	Target prec.	Thr.
CTU-IoT-23						
ddos_high_rate	0.23	DDoS	0.76	DDoS	0.76 (0.00)	0.00
cc_heartbeat_beaconing	0.30	C&C-HB	0.66	C&C-HB	0.66 (0.00)	0.06
benign_single_probe	0.58	Benign	0.38	Benign	0.38 (0.00)	0.00
okiru_low_rate	0.63	Benign	0.35	Okiru	0.31 (0.00)	0.00
asymmetric_flood	0.52	DDoS	0.44	DDoS	0.43 (0.10)	0.09
persistent_connection	0.15	C&C-HB	0.63	C&C-HB	0.63 (0.00)	0.02
incomplete_handshake	0.38	C&C-HB	0.63	C&C-HB	0.63 (0.16)	0.07
ddos_short_burst	0.23	DDoS	0.76	DDoS	0.76 (0.00)	0.00
CIC-IoT-2023						
syn_flood_ddos	0.20	DDoS-SYN	1.00	DDoS-SYN	1.00 (0.00)	0.48
udp_mirai_flood	0.09	Mirai	0.85	Mirai	0.85 (0.14)	0.20
recon_port_scan	0.04	Recon	0.74	Recon	0.74 (0.00)	0.12
vuln_scan	0.00	DDoS-SYN	0.67	VulnScan	0.01 (0.00)	0.28
benign_low_rate	0.70	Mirai	0.28	Benign	0.28 (0.02)	2.90
large_payload_flood	0.06	Benign	0.59	Benign	0.59 (0.06)	1.91
high_entropy_scan	0.00	Recon	0.46	Recon	0.46 (0.05)	0.42
short_connection_flood	0.29	DDoS-SYN	0.63	DDoS-SYN	0.63 (0.09)	0.55

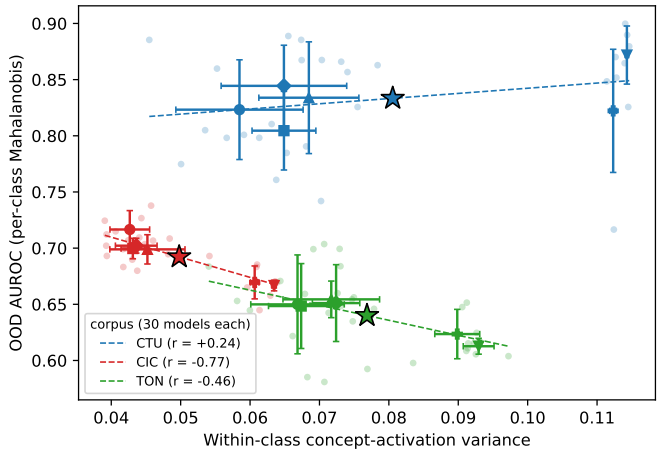


Fig. 4: Within-class concept-activation variance vs. OOD AUROC for the CBM concept-space models, on leak-free feature lists (CTU: known-class-only selection, Section SVII; CIC/TON: fixed). Colour: dataset; marker: configuration; faint points: seeds; solid points: configuration means (±1 seed std); dashed: per-dataset fits; stars: dataset means. Slopes are negative on CIC and TON; on CTU the weakly positive slope is set by SequentialCBM (top right), and the dataset means are non-monotone.

D. Known-Class Geometry: Consistent on Two Corpora, Configuration-Driven on CTU

Across the ten NeSy-NIDS models, mean within-class rule-activation variance tracks OOD AUROC (r=0.971), but the models form two dataset clusters, so this correlation cannot separate a property of the representation from a property of the dataset. We therefore test known-class statistics across the thirty CBM concept-space models per dataset (six configurations × five seeds) on three corpora.

Within-class variance. On CIC-IoT-2023 and TON_IoT more within-class variance goes with worse detection ($r=-0.77$ and -0.46). On CTU-IoT-23 the slope is weakly *positive* ($+0.24$) and comes from SequentialCBM, which shares the highest within-class variance with HybridCBM but is also the strongest CTU concept-space detector (0.872, against 0.822 for HybridCBM; Fig. 4). Without SequentialCBM the CTU correlation is $+0.01$. The within-dataset direction also does not carry across corpora: the three dataset means are non-monotone ($r=+0.36$ over three points), so a corpus's mean known-class variance does not rank its detectability.

TABLE XIII: Within-dataset Pearson r between each known-class statistic and OOD AUROC over 30 CBM concept-space models per dataset (six configurations × five seeds). All corpora use feature selection blind to the held-out families (CTU: known-class-only selection, Section SVII; CIC/TON: fixed lists). CIC and TON share every sign; the CTU signs are opposite and fall to $|r| \leq 0.22$ without SequentialCBM (Supplementary Table S5).

Known-class statistic	CTU	CIC	TON
Within-class variance	+0.24	−0.77	−0.46
Fisher ratio	−0.18	+0.69	+0.47
Ridge-LDA separability	−0.26	+0.46	+0.52
Silhouette	−0.22	+0.81	+0.51
Nearest-centroid margin	−0.24	+0.65	+0.45
Concept-activation entropy	−0.27	+0.80	+0.40
Closed-set F1	−0.37	+0.53	+0.35
Unknown/known distance ratio†	+0.65	+0.56	+0.72

†Uses the unknown flows, so it is a proxy of AUROC, not an independent predictor.

Seven statistics. We tested seven statistics computed from known classes only: within-class variance, Fisher trace ratio $\mathrm{tr}(S_B)/\mathrm{tr}(S_W)$, ridge-LDA separability $\mathrm{tr}((S_W+\varepsilon I)^{-1}S_B)$, silhouette, nearest-centroid margin, concept-activation entropy, and closed-set F1. On CIC-IoT-2023 and TON_IoT each keeps the same sign (Table XIII): more known-class separability, or less within-class variance, goes with higher OOD AUROC, as intuition and Vaze et al. [20] predict, and at the configuration-mean level the six separability and variance statistics reach $|r_6| \geq 0.93$ on both. On CTU-IoT-23 all seven take the opposite sign at the model level ($|r| \leq 0.37$), and leaving out SequentialCBM shrinks them to $|r| \leq 0.22$, with three of them changing sign (Supplementary Table S5). SequentialCBM pairs the lowest closed-set F1 and the least separable concept space with one of the highest CTU AUROCs, so this one configuration sets the CTU direction. The leave-one-configuration-out check shows that the CTU reversal is a property of one configuration rather than of the corpus. Correlations pooled over configurations, as here, therefore need such a check before they are read as a property of a dataset. These are within-corpus associations with Mahalanobis AUROC, not a test that a statistic predicts AUROC on an unseen corpus or selects a scorer; the one predictive procedure we tried (leave-family-out) did not (Section IV-B).

A manipulation and an interpretation. Retraining NeSy-NIDS at annealing steepness $k_T \in \{2, 5, 10, 20\}$ (five seeds, constant known-class F1) raises within-class variance on CIC (0.024→0.037) but AUROC follows only weakly (r=0.32, 0.63 at $k_T \leq 5$ vs. 0.66 at $k_T \geq 10$); on CTU variance barely

moves (0.125–0.128) and is not positively associated with AUROC ($r = -0.15$). A structural argument is consistent with this: a known-class-conditional scorer separates unknowns by where they fall relative to the known class model, which depends on the unknown flows as well as the representation, so a known-class statistic alone need not determine it. We state the result narrowly, for the settings tested.

V. Discussion

Deployment implications. Rejection-threshold calibration and scorer selection are separate problems: the threshold τ needs only known validation traffic and held the known rejection rate near its target (Section IV-B), whereas choosing the scorer does not reduce to known traffic in the same way (limitation (iii)). For model tuning, γ offers no reliable OOD benefit, so effort is better spent matching the scorer to the representation and reading the concept or rule vector alongside the fidelity limits of Section IV-C.

Computational cost. On CTU-IoT-23, JointCBM adds 296 parameters to the MLP and NeSy-NIDS has 13,313, all under 0.4 MB in `float32`. Per-flow (batch-1) latency on one thread of an Intel Core i9-9820X is 96 μs (MLP), 109 μs (JointCBM) and about 590 μs (NeSy-NIDS, an un-vectorised per-rule loop), measured with the submitted 37-feature input.

Limitations. (i) *Authored vocabulary.* The concepts and rules were written by us from documented attack signatures [1], [2] and not validated by operators or independent experts. (ii) *Capture overlap and feature selection.* Splits are random at the flow level; capture-, session- and device-disjoint evaluation remains untested, and the few source hosts per CTU family (4–7) make it harder to construct. The CTU feature set was selected with the held-out labels visible; re-selecting without them per split leaves the findings in place. Each split is a single class-balanced sample, so multi-sample or full-corpus evaluation is a natural next step (Sections III-A, SVII). (iii) *Scorer selection.* No procedure we tested selects the best scorer from known traffic alone (Section IV-B). (iv) *Timing* was measured only on a desktop CPU. (v) *Coverage.* The post-hoc detector comparison uses the seed-0 MLP and the (λ, γ) grid three seeds; NeSy-NIDS and the temporal/alternative splits cover two datasets, the geometry three. All corpora are fixed-vector flow representations with hand-authored concepts, so we make no claim for packet-sequence, graph, or learned-concept representations.

VI. Conclusion

Interpretable IoT intrusion detectors can match an unconstrained MLP on known classes, but their inspectable concepts and rules do not account for their decisions, and their open-set detection depends on dataset, architecture and scorer; the known-class statistics we measured are associated with detectability on two corpora, through a single configuration on the third, and none is validated as a predictor. Evaluations of interpretable NIDS should report decision fidelity alongside accuracy, support open-set claims with multiple seeds, architectures, datasets and leak-free feature selection, and check pooled correlations with a leave-one-configuration-out analysis; the most important remaining test is a capture-disjoint evaluation.

References

[1] Y. Mirsky, T. Doitshman, Y. Elovici, and A. Shabtai, "Kitsune: An ensemble of autoencoders for online network intrusion detection," in *Network and Distributed System Security Symposium (NDSS)*, 2018.

[2] Y. Meidan *et al.*, "N-BaIoT: Network-based detection of IoT botnet attacks using deep autoencoders," *IEEE Pervasive Computing*, vol. 17, no. 3, pp. 12–22, 2018.

[3] C. Rudin, "Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead," *Nature Machine Intelligence*, vol. 1, no. 5, pp. 206–215, 2019.

[4] D. Arp *et al.*, "Dos and don'ts of machine learning in computer security," in *Proceedings of the 31st USENIX Security Symposium*. USENIX Association, 2022, pp. 3971–3988.

[5] F. Wei, H. Li, Z. Zhao, and H. Hu, "xNIDS: Explaining deep learning-based network intrusion detection systems for active intrusion responses," in *Proceedings of the 32nd USENIX Security Symposium*. USENIX Association, 2023, pp. 4337–4354.

[6] D. Han *et al.*, "Anomaly detection in the open world: Normality shift detection, explanation, and adaptation," in *Proceedings of the 2023 Network and Distributed System Security Symposium (NDSS)*. Internet Society, 2023.

[7] P. W. Koh *et al.*, "Concept bottleneck models," in *Proceedings of the 37th International Conference on Machine Learning (ICML)*. PMLR, 2020, pp. 5338–5348.

[8] F. Petersen, C. Borgelt, H. Kuehne, and O. Deussen, "Deep differentiable logic gate networks," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 35, 2022.

[9] Z. Yang, J. Lee, and C. Park, "Injecting logical constraints into neural networks via straight-through estimators," in *Proceedings of the 39th International Conference on Machine Learning (ICML)*. PMLR, 2022, pp. 25 096–25 122.

[10] K. Lee, K. Lee, H. Lee, and J. Shin, "A simple unified framework for detecting out-of-distribution samples and adversarial attacks," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 31, 2018.

[11] D. Han *et al.*, "Rules refine the riddle: Global explanation for deep learning-based anomaly detection in security applications," in *Proceedings of the 2024 ACM SIGSAC Conference on Computer and Communications Security (CCS)*. ACM, 2024.

[12] M. E. Zarlenga *et al.*, "Concept embedding models: Beyond the accuracy-explainability trade-off," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 35, 2022.

[13] M. Yuksekgonul, M. Wang, and J. Zou, "Post-hoc concept bottleneck models," in *Proceedings of the 11th International Conference on Learning Representations (ICLR)*, 2023.

[14] A. Mahinpei, J. Clark, I. Lage, F. Doshi-Velez, and W. Pan, "Promises and pitfalls of black-box concept learning models," *arXiv preprint arXiv:2106.13314*, 2021.

[15] S. Shoukat, T. Gao, D. Javeed, and M. Adil, "From evidence to decision: Concept–prototype reasoning for intrusion detection in SDN-based industrial networks," *Computer Networks*, 2026, in press.

[16] Y. Bengio, N. Léonard, and A. Courville, "Estimating or propagating gradients through stochastic neurons for conditional computation," *arXiv preprint arXiv:1308.3432*, 2013.

[17] S. Fort, J. Hu, and B. Lakshminarayanan, "Exploring the limits of out-of-distribution detection," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 34, 2021.

[18] D. Hendrycks and K. Gimpel, "A baseline for detecting misclassified and out-of-distribution examples in neural networks," in *Proceedings of the 5th International Conference on Learning Representations (ICLR)*, 2017.

[19] W. Liu, X. Wang, J. D. Owens, and Y. Li, "Energy-based out-of-distribution detection," in *Advances in Neural Information Processing Systems*, vol. 33, 2020.

[20] S. Vaze, K. Han, A. Vedaldi, and A. Zisserman, "Open-set recognition: A good closed-set classifier is all you need?" in *International Conference on Learning Representations (ICLR)*, 2022.

[21] L. Yang *et al.*, "CADE: Detecting and explaining concept drift samples for security applications," in *Proceedings of the 30th USENIX Security Symposium*. USENIX Association, 2021, pp. 3237–3254.

[22] G. Apruzzese *et al.*, "The role of machine learning in cybersecurity," *Digital Threats: Research and Practice*, vol. 4, no. 1, 2023.

[23] F. De Keersmaeker, Y. Cao, G. K. Ndonda, and R. Sadre, "A survey of public IoT datasets for network security research," *IEEE Communications Surveys & Tutorials*, vol. 25, no. 3, pp. 1808–1840, 2023.

[24] A. Parmisano, S. Garcia, and M. J. Erquiaga, "CTU-IoT-23: A labeled dataset with malicious and benign IoT network traffic," *Stratosphere Laboratory Technical Report*, 2020, czech Technical University in Prague.
[25] E. C. P. Neto, S. Dadkhah, R. Ferreira, A. Zohourian, R. Lu, and A. A. Ghorbani, "CICIoT2023: A real-time dataset and benchmark for large-scale attacks in IoT environment," *Sensors*, vol. 23, no. 13, p. 5941, 2023.
[26] A. Alsaedi, N. Moustafa, Z. Tari, A. Mahmood, and A. Anwar, "TON_IoT telemetry dataset: A new generation dataset of IoT and IIoT for data-driven intrusion detection systems," *IEEE Access*, vol. 8, pp. 165 130–165 150, 2020.
[27] L. Breiman, "Random forests," *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.
[28] Y. Sun, Y. Ming, X. Zhu, and Y. Li, "Out-of-distribution detection with deep nearest neighbors," in *International Conference on Machine Learning (ICML)*, 2022.

Ngari Crisphine Macharia (Student Member, IEEE) is pursuing the M.Sc. degree at Southern Illinois University Carbondale, USA (B.Tech., GITAM, India, 2024), working on machine learning for network security.

Ning Weng (Senior Member, IEEE) is a Professor of Electrical, Computer and Biomedical Engineering at Southern Illinois University Carbondale (Ph.D., UMass Amherst, 2005), working on deep packet inspection and network security.